

Bayesian Quantile Regression using a Mixture of Pólya Tree Prior

April 5, 2014

Abstract

In this paper, we develop Bayesian methods for quantile regression. We adopt a mixture of Pólya tree prior for the regression error term and account for the change of quantile regression parameter via heterogeneity of the error term. By marginalizing the predictive density function of the Pólya tree distribution, quantile regression parameters can be obtained in closed form and exact inference can be made through MCMC. Our method allow simultaneous quantile regressions and avoids the problem of crossing quantile lines. The model is applied to data from a clinical trial on weight management.

1 Introduction

Quantile regression is a powerful way of studying the relationship between response and covariates when one (or several) quantiles are of interest. Quantile regression offers a more complete description of the conditional distribution of the response. The dependence between upper or lower quantiles of the response variable and the covariates typically vary differentially from the mean. The dependence of quantiles on covariates is often of interest in econometrics, educational studies, biomedical studies, and environment studies (Yu and Moyeed, 2001; Buchinsky, 1994, 1998; He et al., 1998; Koenker and Machado, 1999; Wei et al., 2006; Yu et al., 2003). A comprehensive review of quantile regression was presented by Koenker (2005).

The traditional frequentist approach was proposed by Koenker and Bassett (1978) for a single quantile (τ) with estimators derived by minimizing a loss check function $\sum_{i=1}^n \rho_\tau(y_i - x_i' \beta)$, where $\rho_\tau(\epsilon) = \epsilon(\tau - I(\epsilon < 0))$. They do not make any distributional assumptions for residuals and use linear programming techniques for estimation. The popularity of this approach is due to its computational efficiency, well-developed asymptotic properties, and straightforward extensions to simultaneous quantile regression and random effect models. However, asymptotic inference may not be accurate for small sample sizes.

Bayesian approaches offer exact inference. Motivated by the loss check function, Yu and Moyeed (2001) proposed an asymmetric Laplace distribution for the error term, such that maximizing the posterior distribution is equivalent to minimizing the check function. Other than parametric Bayesian approaches, some semiparametric methods have been proposed for *median regression*. Walker and Mallick (1999) used a diffuse finite Pólya Tree prior for the

error term. [Kottas and Gelfand \(2001\)](#) modeled the error by two families of median zero distribution using a mixture Dirichlet process priors, which is very useful for unimodal error distributions. [Hanson and Johnson \(2002\)](#) adopted mixture of Pólya Tree prior on error term to make inference in regression model. They illustrated the implementation on AFT model for the median survival time, which showed robustness of Pólya in terms of multimodality and skewness. [Reich et al. \(2010\)](#) used an infinite mixture of Gaussian densities on the residual. Other recent approaches include quantile pyramid priors, mixture of Dirichlet process priors of multivariate normal distributions and infinite mixture of Gaussian densities which put quantile constraints on the residuals ([Hjort and Petrone, 2007](#); [Hjort and Walker, 2009](#); [Kottas and Krnjajić, 2009](#)).

Like the asymmetric Laplace distribution, all of the above methods are single semiparametric quantile regression methods, which have some limitations. The densities have their restrictive mode at the quantile of interest, which is not appropriate when extreme quantiles are being investigated. Other criticisms include crossed quantile lines, monotonicity constraints and difficulty in making inference for quantile regression parameter for an interval of τ s. Joint inference is poor in borrowing information through single quantile regressions. It is not coherent to pool from every individual quantile regression, because the sampling distribution of Y for τ_1 is usually different from that under quantile τ_2 since they are assuming different error distribution under two different quantile regressions.

In order to solve those problems, simultaneous linear quantile regression have been proposed by [Tokdar and Kadane \(2011\)](#), in which an interpolation of two monotone curves was introduced to put constraints on monotonicity of quantile lines. Another popular approach is to assign a nonparametric model for the error term to avoid the monotonicity problem. ([Scaccia and Green, 2003](#)) used mixture of normal distributions with an unknown number of components to model the distribution of measurement by B-splines. ([Geweke and Keane, 2007](#)) extends the traditional Bayesian mixture of normals model assuming state probabilities depending on observed covariates. ([Taddy and Kottas, 2010](#)) considered the model with joint distribution of responses and covariates using flexible Dirichlet process mixture models.

We use a mixture of Pólya Tree (PT) priors in our approach. PT priors were introduced decades ago ([Freedman, 1963](#); [Fabius, 1964](#); [Ferguson, 1974](#)) and [Lavine \(1992, 1994\)](#) extended them to Pólya Tree models. The major advantage of Pólya Tree over Dirichlet process is that it can be absolutely continuous with probability 1 and it can be easily tractable. In a regression context, [Walker and Mallick \(1997, 1999\)](#) assigned a finite Pólya Tree prior to the random effects in a generalized linear mixed model. [Berger and Guglielmi \(2001\)](#) used a mixture of Pólya Tree comparing data distribution coming from parametric distribution or mixture of Pólya Tree. They used a Pólya tree process to test the fit of data to a parametric model by embedding the parametric model in a nonparametric alternative and computing the Bayes factor of the parametric model to the nonparametric alternative. As mentioned earlier [Hanson and Johnson \(2002\)](#) modeled the error term as a mixture of Pólya tree prior in the regression model.

In this article, we present a Bayesian approach by adopting a mixture of Pólya Tree prior for the regression error term, and we account for the change of quantile regression parameter via heterogeneity of the error term. As a result, several quantile regression can be fit simultaneously and there is a closed form for posterior quantile regression parameter. Exact inference can be made through Monte Carlo Markov Chain (MCMC) approach, and our method avoids the problem of crossing quantile lines that occurs in the traditional frequentist

quantile regressions.

The rest of the paper is organized as follows. In section 2, we introduce the heterogeneity model and derive a closed form for marginalized posterior quantile regression parameter with mixture of Pólya tree prior. In section 3, we conduct some simulation studies and apply our approach on a real data example in section 4. Finally, conclusions and discussions are presented in section 5.

2 Model, Priors, and Computations

2.1 Model

Let Y be a random variable with CDF F . The τ th quantile of Y is defined as

$$Q_Y(\tau) = \inf_y \{y : F(y) \geq \tau\}.$$

If covariates x_1, \dots, x_n are of interest, then the quantile regression parameter satisfies the following condition:

$$Q_Y(\tau) = X' \beta(\tau),$$

where X is the matrix of covariates including an intercept. If F is continuous, then $F(X' \beta(\tau)) = \tau$, i.e., $p(Y \leq X' \beta(\tau)) = \tau$.

Now, consider a location shift model,

$$y_i = x_i \beta + \epsilon_i,$$

where $\epsilon_i \stackrel{\text{i.i.d}}{\sim} F_\epsilon$. Then, the τ th quantile regression parameter can be expressed as

$$\beta(\tau) = \beta + F_\epsilon^{-1}(\tau) e_1, \tag{1}$$

where $e_1 = [1, 0, \dots, 0]^T$, and $F_\epsilon^{-1}(\tau)$ is the τ th quantile for error ϵ .

As we can see from equation (1), if the model is homogeneous, i.e., i.i.d case, then for different quantiles τ , the corresponding quantile regression parameters only vary in the first component, the intercept. The rest of the quantile regression parameters stay the same. Therefore, quantile lines for different quantiles are parallel to each other.

Now, consider the heterogeneous linear regression model from He et al. (1998)

$$y_i = x_i' \beta + (x_i' \gamma) \epsilon_i, \tag{2}$$

where $x_i' \gamma$ is positive for all i . Under this model, the τ th quantile regression parameter is

$$\beta(\tau) = \beta + F_\epsilon^{-1}(\tau) \gamma, \tag{3}$$

Quantile lines are no longer parallel under the heterogeneous linear model which adds considerably more flexibility.

We use a mixture of Pólya Tree prior for the error term in equation (2) and derive a closed form for posterior quantile regression parameter in (3). Since a Pólya tree is a very flexible way to model the unknown distribution, our approach makes fewer assumptions. Exact inference can be made through MCMC using functionals of posterior samples. The next subsection briefly reviews the Pólya tree priors and their relevant properties.

2.2 Pólya Tree

Lavine (1992, 1994) and Mauldin et al. (1992) developed theory for Pólya trees priors as a generalization of the Dirichlet process (Ferguson, 1974). Denote $E = \{0, 1\}$ and E^m as the m -fold product of E , $E^0 = \emptyset$, $E^* = \cup_0^\infty E^m$ and Ω be a separable measurable space, $\pi_0 = \Omega$, $\Pi = \{\pi_m : m = 0, 1, \dots\}$ be a separating binary tree of partitions of Ω . In addition, define $B_\emptyset = \Omega$ and $\forall \epsilon = \epsilon_1 \cdots \epsilon_m \in E^*$, $B_{\epsilon 0}$ and $B_{\epsilon 1}$ are the two partition of B_ϵ .

Definition 2.1. A random probability measure G on (Ω, \mathcal{F}) is said to have a Pólya tree distribution, or a Pólya tree prior with parameter (Π, \mathcal{A}) , written as $G|\Pi, \mathcal{A} \sim \text{PT}(\Pi, \mathcal{A})$, if there exist nonnegative numbers $\mathcal{A} = \{\alpha_\epsilon, \epsilon \in E^*\}$ and random vectors $\mathcal{Y} = \{Y_\epsilon : \epsilon \in E^*\}$ such that the following hold:

1. all the random variables in \mathcal{Y} are independent;
2. $Y_\epsilon = (Y_{\epsilon 0}, Y_{\epsilon 1}) \sim \text{Dirichlet}(\alpha_{\epsilon 0}, \alpha_{\epsilon 1}), \forall \epsilon \in E^*$;
3. $\forall m = 1, 2, \dots$, and $\forall \epsilon \in E^*$, $G(B_{\epsilon_1, \dots, \epsilon_m}) = \prod_{j=1}^m Y_{\epsilon_1 \dots \epsilon_j}$.

2.2.1 Pólya Tree Parameters

There are two parameters in the Pólya tree distribution (Π, \mathcal{A}) . A Pólya tree is centered around a pre-specified distribution G_0 , which is called the baseline measure. The \mathcal{A} family determines how much G can deviate from G_0 . Ferguson (1974) pointed out $\alpha_\epsilon = 1$ yields a G that is absolutely continuous with probability 1, and $\alpha_{\epsilon_1, \dots, \epsilon_m} = m^2$ yields G that is absolutely continuous with probability 1. Walker and Mallick (1999) and Paddock (1999) considered $\alpha_{\epsilon_1, \dots, \epsilon_m} = cm^2$, where $c > 0$. Berger and Guglielmi (2001) considered $\alpha_{\epsilon_1, \dots, \epsilon_m} = c\rho(m)$. In general, any $\rho(m)$ such that $\sum_{m=1}^\infty \rho(m)^{-1} < \infty$ guarantees G to be absolutely continuous. In our case, we adopt $\alpha_{\epsilon_1, \dots, \epsilon_m} = cm^2$.

As to the partition parameter Π , the canonical way of constructing a Pólya tree distribution G centering on G_0 , a continuous CDF is to choose $B_0 = G_0^{-1}([0, 1/2])$, $B_1 = G_0^{-1}((1/2, 1])$, such that $G(B_0) = G(B_1) = 1/2$. Furthermore, for all $\epsilon \in E^*$, choose $B_{\epsilon 0}$ and $B_{\epsilon 1}$ to satisfy $G(B_{\epsilon 0}|B_\epsilon) = G(B_{\epsilon 1}|B_\epsilon) = 1/2$, then any choice of \mathcal{A} makes G coincide with G_0 . A simple example is to choose $B_{\epsilon 0}$ and $B_{\epsilon 1}$ in level m by setting them as $G_0^{-1}((k/2^m, (k+1)/2^m])$, for $k = 0, \dots, 2^m - 1$.

2.2.2 Some properties of Pólya Tree

Suppose $G \sim \text{PT}(\Pi, \mathcal{A})$ is a random probability measure and $\epsilon_1, \epsilon_2, \dots$ are random samples from G .

Definition 2.2 (Expectation of Pólya Tree). $F = E(G)$ as a probability measure is defined by $F(B) = E(G(B)), \forall B \in \mathcal{B}$. By the definition of Pólya tree, for any $\epsilon \in E^*$,

$$F(B_\epsilon) = E(G(B_\epsilon)) = \prod_{j=1}^m \frac{\alpha_{\epsilon_1, \dots, \epsilon_j}}{\alpha_{\epsilon_1, \dots, \epsilon_{j-1}, 0} + \alpha_{\epsilon_1, \dots, \epsilon_{j-1}, 1}}.$$

Remark 2.3. If G is constructed based on baseline measure G_0 and we set $\alpha_{\epsilon_1, \dots, \epsilon_m} = cm^2$, $\alpha_{\epsilon 0} = \alpha_{\epsilon 1}$, then $\forall B \in \mathcal{B}, F(B) = G_0(B)$; thus, $F = G_0$, if there is no data.

Definition 2.4 (Density Function). Suppose $F = E(G), G|\Pi, \mathcal{A} \sim \text{PT}(\Pi, \mathcal{A})$, where G_0 is the baseline measure. Then, using the canonical construction, $F = G_0$ (as shown above), the density function is

$$f(y) = \left[\prod_{j=1}^m \frac{\alpha_{\epsilon_1, \dots, \epsilon_j}(y)}{\alpha_{\epsilon_1, \dots, \epsilon_{j-1}, 0}(y) + \alpha_{\epsilon_1, \dots, \epsilon_{j-1}, 1}(y)} \right] 2^m g_0(y), \quad (4)$$

where g_0 is the pdf of G_0 .

Remark 2.5. When using the canonical construction with no data, $\alpha_{\epsilon_0} = \alpha_{\epsilon_1}$, equation (4) simplifies to

$$f(y) = g_0(y).$$

Remark 2.6 (Conjugacy). If $y_1, \dots, y_n | G \sim G, G|\Pi, \mathcal{A} \sim \text{PT}(\Pi, \mathcal{A})$, then $G|y_1, \dots, y_n, \Pi, \mathcal{A} \sim \text{PT}(\Pi, \mathcal{A}^*)$, where in $\mathcal{A}^*, \forall \epsilon \in E^*$,

$$\alpha_{\epsilon}^* = \alpha_{\epsilon} + n_{\epsilon}(y_1, \dots, y_n),$$

where $n_{\epsilon}(y_1, \dots, y_n)$ indicates the count of how many samples of y_1, \dots, y_n fall in B_{ϵ} .

2.2.3 Mixture of Pólya Trees

The behavior of a single Pólya tree highly depends on how the partition is specified. A random probability measure G_{θ} is said to be a mixture of Pólya tree if there exists a random variable θ with distribution h_{θ} , and Pólya tree parameters $(\Pi^{\theta}, \mathcal{A}^{\theta})$ such that $G_{\theta}|\theta = \theta \sim \text{PT}(\Pi^{\theta}, \mathcal{A}^{\theta})$.

Example 2.7. Suppose $G_0 = N(\mu, \sigma^2)$ is the baseline measure. For $\epsilon \in E^*, \alpha_{\epsilon_m} = cm^2, \theta = (\mu, \sigma, c)$ is the mixing index and the distribution on $\Theta = (\mu, \sigma, c)$ is the mixing distribution.

With the mixture of Pólya tree, the influence of the partition is lessened. Thus, inference will not be affected greatly by a single Pólya tree distribution.

2.2.4 Predictive Error Density, Cumulative Density Function and Quantiles

Suppose G_{θ} is the baseline measure, $g_0(y)$ is the density function. Π^{θ} is defined as

$$B_{\epsilon_1, \dots, \epsilon_m}^{\theta} = \left(G_{\theta}^{-1} \left(\frac{k}{2^m} \right), G_{\theta}^{-1} \left(\frac{k+1}{2^m} \right) \right),$$

where k is the index of partition $\epsilon_1, \dots, \epsilon_m$ in level m . \mathcal{A}^c is defined as

$$\alpha_{\epsilon_1, \dots, \epsilon_m} = cm^2.$$

Therefore, the error model is

$$y_1, \dots, y_n | G_{\theta} \stackrel{\text{i.i.d}}{\sim} G, \\ G | \Pi^{\theta}, \mathcal{A}^c \sim \text{PT}(\Pi^{\theta}, \mathcal{A}^c).$$

The predictive density function of $Y|y_1, \dots, y_n, \theta$, marginalizing out G , is

$$f_Y^\theta(y|y_1, \dots, y_n) = \lim_{m \rightarrow \infty} \left(\prod_{j=2}^m \frac{cj^2 + n_{\epsilon_1 \dots \epsilon_j(y)}(y_1, \dots, y_n)}{2cj^2 + n_{\epsilon_1 \dots \epsilon_{j-1}(y)}(y_1, \dots, y_n)} \right) 2^{m-1} g_0(y), \quad (5)$$

where $n_{\epsilon_1 \dots \epsilon_j(y)}(y_1, \dots, y_n)$ denotes the number of observations y_1, \dots, y_n dropping in the bin $\epsilon_1 \dots \epsilon_j$ where y stays in the level j . Notice that, if we restrict the first level weight as $\alpha_0 = \alpha_1 = 1$, then we only need to update levels beyond the first level.

Remark 2.8 (The predictive density for Finite Pólya Tree). *In practice, a finite M level Pólya Tree is usually adopted to approximate the full Pólya tree, in which, only up to M levels are updated. The corresponding predictive density becomes*

$$f_Y^{\theta, M}(y|y_1, \dots, y_n) = \left(\prod_{j=2}^M \frac{cj^2 + n_{\epsilon_1 \dots \epsilon_j(y)}(y_1, \dots, y_n)}{2cj^2 + n_{\epsilon_1 \dots \epsilon_{j-1}(y)}(y_1, \dots, y_n)} \right) 2^{M-1} g_0(y). \quad (6)$$

The rule of thumb for choosing M is to set $M = \log_2 n$, where n is the sample size ([Hanson and Johnson, 2002](#)).

[Hanson and Johnson \(2002\)](#) showed the approximation to (5) given in (6) is exact for M large enough. We now derive the predictive cdf and the predictive quantile(s).

Theorem 2.9. *Based on the predictive density function (6) of a finite Pólya tree distribution, the predictive cumulative density function is*

$$F_Y^{\theta, M}(y|y_1, \dots, y_n) = \sum_{i=1}^{N-1} P_i + P_N \left(G_\theta(y) 2^M - (N-1) \right), \quad (7)$$

where

$$P_i = \frac{1}{2} \left(\prod_{j=2}^M \frac{cj^2 + n_{j, \lceil i2^{j-M} \rceil}(y_1, \dots, y_n)}{2cj^2 + n_{j-1, \lceil i2^{j-1-M} \rceil}(y_1, \dots, y_n)} \right) \text{ and} \\ N = \left\lceil 2^M G_\theta(y) + 1 \right\rceil,$$

in which $n_{j, \lceil i2^{j-M} \rceil}(y_1, \dots, y_n)$ denotes the number of observations y_1, \dots, y_n in the $\lceil i2^{j-M} \rceil$ slot at level j , $\lceil \cdot \rceil$ is the ceiling function, and $\lfloor \cdot \rfloor$ is the floor function.

Proof.

$$\begin{aligned}
F_Y^{\theta,M}(y|y_1, \dots, y_n) &= \int_{-\infty}^y f_Y^{\theta,M}(y|y_1, \dots, y_n) dx \\
&= \int_{-\infty}^y \left(\prod_{j=2}^M \frac{cj^2 + n_{\epsilon_1 \dots \epsilon_j(y)}(y_1, \dots, y_n)}{2cj^2 + n_{\epsilon_1 \dots \epsilon_{j-1}(y)}(y_1, \dots, y_n)} \right) 2^{M-1} g_\theta(y) dy \\
&= \sum_{i=1}^{N-1} \left(\prod_{j=2}^M \frac{cj^2 + n_{j, \lceil i2^{j-1-M} \rceil}(y_1, \dots, y_n)}{2cj^2 + n_{j-1, \lceil i2^{j-1-M} \rceil}(y_1, \dots, y_n)} 2^{M-1} \int_{\epsilon_{M,i}} g_\theta(y) dy \right) \\
&\quad + \int_{G_\theta^{-1}((N-1)/2^M)}^y \left(\prod_{j=2}^M \frac{cj^2 + n_{j, \lceil N2^{j-1-M} \rceil}(y_1, \dots, y_n)}{2cj^2 + n_{j-1, \lceil N2^{j-1-M} \rceil}(y_1, \dots, y_n)} \right) 2^{M-1} g_\theta(y) dy \\
&= \sum_{i=1}^{N-1} P_i + P_N 2^M \left(G_\theta(y) - G_\theta(G_\theta^{-1} \left(\frac{N-1}{2^M} \right)) \right) \\
&= \sum_{i=1}^{N-1} P_i + P_N \left(G_\theta(y) 2^M - (N-1) \right),
\end{aligned}$$

where $\epsilon_{M,i}$ is the i th partition in level M . □

Theorem 2.10. *The posterior predictive quantile of finite Pólya tree distribution is*

$$Q_{Y|y_1, \dots, y_n}^{\theta,M}(\tau) = G_\theta^{-1} \left(\frac{\tau - \sum_{i=1}^N P_i + NP_N}{2^M P_N} \right), \quad (8)$$

where N satisfies $\sum_{i=1}^{N-1} P_i < \tau \leq \sum_{i=1}^N P_i$.

Proof. From equation (7),

$$\begin{aligned}
\tau = F_Y^{\theta,M}(y|y_1, \dots, y_n) &= \sum_{i=1}^{N-1} P_i + P_N \left(G_\theta(y) 2^M - (N-1) \right) \\
\Rightarrow G_\theta(y) &= \frac{\tau - \sum_{i=1}^N P_i + NP_N}{2^M P_N} \\
y &= G_\theta^{-1} \left(\frac{\tau - \sum_{i=1}^N P_i + NP_N}{2^M P_N} \right).
\end{aligned}$$
□

Now the explicit form for quantile regression coefficients in equation (3) becomes:

$$\beta(\tau) = \beta + \gamma G_\theta^{-1} \left(\frac{\tau - \sum_{i=1}^N P_i + NP_N}{2^M P_N} \right), \quad (9)$$

where P_i and N are the notations in equation (7) and (8). This will greatly facilitate computations.

2.3 Fully Bayesian Quantile Regression Specification with Mixture of Pólya Tree Priors

The full Bayesian specification of quantile regression is given as follows,

$$\begin{aligned} y_i &= \mathbf{x}_i' \boldsymbol{\beta} + (\mathbf{x}_i' \boldsymbol{\gamma}) \epsilon_i, i = 1, \dots, n \\ \epsilon_i | G_\theta &\stackrel{\text{i.i.d}}{\sim} G_\theta \\ G_\theta | \Pi^\theta, \mathcal{A}^\theta &\sim \text{PT}(\Pi^\theta, \mathcal{A}^\theta) \\ \boldsymbol{\theta} = (\sigma, c) &\sim \pi_\theta(\boldsymbol{\theta}) \\ \boldsymbol{\beta} &\sim \pi_\beta(\boldsymbol{\beta}) \\ \boldsymbol{\gamma} &\sim \pi_\gamma(\boldsymbol{\gamma}). \end{aligned}$$

In order to not confound the location parameter, ϵ_i or G is set to have median 0 by fixing $\alpha_0 = \alpha_1 = 1$. For the similar reason, the first component of $\boldsymbol{\gamma}$ is fixed at 1.

As to the monotonicity of quantile lines, $\mathbf{x}_i' \boldsymbol{\gamma} > 0$ is required for all \mathbf{x}_i . We introduce a latent variable approach borrowing idea in [Reich et al. \(2011\)](#); [Reich \(2012\)](#); [Reich and Smith \(2013\)](#). We scale $X_{ij} \in [-1, 1]$, therefore $\mathbf{x}_i' \boldsymbol{\gamma}$ is minimized as $1 - \sum_{i=2}^p |\gamma_i|$ for the worst case ("WC"). To meet the constraints for all \mathbf{x}_i , we introduce latent coefficients $\boldsymbol{\gamma}^* = (1, \gamma_2^*, \dots, \gamma_p^*)$ such that

$$\gamma_i = \begin{cases} \gamma_i^*, & \text{WC}(\boldsymbol{\gamma}^*) > 0 \\ 0, & \text{WC}(\boldsymbol{\gamma}^*) < 0 \end{cases}.$$

[Reich and Smith \(2013\)](#) showed the restrictions on $\boldsymbol{\gamma}$ also provide a very flexible model.

The posterior distribution of $(\boldsymbol{\beta}, \boldsymbol{\gamma}^*, \sigma, c)$ is given as

$$\begin{aligned} p(\boldsymbol{\beta}, \boldsymbol{\gamma}^*, \sigma, c | \mathbf{Y}) &\propto L(\mathbf{Y} | \boldsymbol{\beta}, \boldsymbol{\gamma}^*, \sigma, c) \pi_\beta(\boldsymbol{\beta}) \pi_{\boldsymbol{\gamma}^*}(\boldsymbol{\gamma}^*) \pi_\sigma(\sigma) \pi_c(c) \\ &= \frac{1}{\prod_{i=1}^n (\mathbf{x}_i' \boldsymbol{\gamma})} p(\epsilon_1, \dots, \epsilon_n | \boldsymbol{\beta}, \boldsymbol{\gamma}, \sigma, c) \pi_\beta(\boldsymbol{\beta}) \pi_{\boldsymbol{\gamma}^*}(\boldsymbol{\gamma}^*) \pi_\sigma(\sigma) \pi_c(c) \\ &= \frac{1}{\prod_{i=1}^n (\mathbf{x}_i' \boldsymbol{\gamma})} p(\epsilon_n | \epsilon_1, \dots, \epsilon_{n-1}, \boldsymbol{\beta}, \boldsymbol{\gamma}, \sigma, c) \cdots p(\epsilon_2 | \epsilon_1, \boldsymbol{\beta}, \boldsymbol{\gamma}, \sigma, c) p(\epsilon_1 | \boldsymbol{\beta}, \boldsymbol{\gamma}, \sigma, c) \\ &\quad \pi_\beta(\boldsymbol{\beta}) \pi_{\boldsymbol{\gamma}^*}(\boldsymbol{\gamma}^*) \pi_\sigma(\sigma) \pi_c(c), \end{aligned} \tag{10}$$

where $\epsilon_i = (y_i - \mathbf{x}_i' \boldsymbol{\beta}) / (\mathbf{x}_i' \boldsymbol{\gamma})$.

For priors of σ and c , we use diffuse gamma distributions,

$$\begin{aligned} \pi(\sigma) &\sim \Gamma(1/2, 1/2), \\ \pi(c) &\sim \Gamma(1/2, 1/2). \end{aligned}$$

We consider two choices for the priors for the parameters $(\boldsymbol{\beta}, \boldsymbol{\gamma}^*)$. First, we consider diffuse p-dimensional normal distributions. Second, we consider spike and slab priors on $(\boldsymbol{\beta}, \boldsymbol{\gamma}^*)$ are an alternative that can do shrinkage and/or variable selection on both quantile regression parameters and heterogeneity parameters and improve efficiency; note that we keep the covariates the same for both to facilitate the form of the regression coefficients in (9). For the latter, we specify continuous spike and slab priors ([George and McCulloch, 1993](#)) on $(\boldsymbol{\beta}, \boldsymbol{\gamma}^*)$.

In particular, the prior for j^{th} component of β , β_j can be written as:

$$\begin{aligned}\pi_{\beta}(\beta_j) &= \delta_{\beta_j} \phi(\beta_j; 0, s_j^2 \sigma_{\beta_j}^2) + (1 - \delta_{\beta_j}) \phi(\beta_j; \beta_j^p, \sigma_{\beta_j}^2), \\ \delta_{\beta_j} &\sim \text{Bernoulli}(\pi_{\beta_j}),\end{aligned}$$

where $\phi(x; \mu, \sigma^2)$ is the density function of normal distribution at x with mean μ and variance σ^2 . $\beta_j^p, \sigma_{\beta_j}^2$ are the mean and variance of the diffuse normal prior for the slab component. δ_{β_j} is the indicator that β_j comes from spike component or from slab component and π_{β_j} is its corresponding probability. $s_j (> 0)$ is small enough so that if $\delta_{\beta_j} = 1$, it indicates $|\beta_j| < 3s_j \sigma_{\beta_j}$ with high probability, thus it can be approximately estimated as 0 and regarded as non-significant and removed from the model; however, we do not formally do this.

We choose β^p , the mean of normal distribution of slab component, to be least square estimates of β given covariates matrix X , i.e., $(X^T X)^{-1} X^T Y$ and $\sigma_{\beta_j}^2$ to be the diagonal component of matrix $\hat{\sigma}^2 (X^T X)^{-1}$, where $\hat{\sigma}^2 = \sum_i^n (y_i - x_i \beta^p)^2 / (n - p)$.

The priors for γ^* are similar to priors for β . But we choose γ^p and σ_{γ} to be **0** and **100**. The parameters π_{β_j} and π_{γ_j} control the belief that the corresponding regressors are needed in the model and we assign Beta(1,1) priors for them.

2.4 Computational Details

In this section, we describe how to draw posterior samples to make inference in our proposed Bayesian quantile regression model with Pólya tree priors using an MCMC algorithm. Functions are written using Fortran within R (R Core Team, 2013) following R library *DPpackage* (Jara et al., 2011). And we have incorporated those functions implementing the algorithm into the new R (R Core Team, 2013) package “bqrpt”.

We use Metropolis-Hasting algorithm to draw posterior samples. The posterior distributions of $(\beta, \gamma^*, \sigma, c | Y)$ are proportional to (10).

We use $\beta_j^c \sim N(\beta_j^{l-1}, t_{\beta_j} (X'X)_{jj}^{-1})$ as candidate distribution for β_j in l -th iteration, where t_{β_j} is the tuning parameter for β_j to adjust acceptance rate (Jara et al., 2009). Similarly, we use $\gamma_j^c \sim N(\gamma_j^{l-1}, t_{\gamma_j} (X'X)_{jj}^{-1})$ as candidate distribution for γ_j^* in l -th iteration. For baseline (centering) normal distribution parameter σ (μ is fixed at 0 due to not confound with location parameter β), we use the lognormal candidate distribution $\sigma^* \sim \text{LN}(\log \sigma^{l-1}, t_{\sigma})$. The same strategy is applied for Pólya tree weight parameter c , where $c^* \sim \text{LN}(\log c^{l-1}, t_c)$.

When using spike and slab priors, $(\delta_{\beta} | \beta, \pi_{\beta})$, $(\delta_{\gamma} | \gamma^*, \pi_{\gamma})$, $(\pi_{\beta} | \delta_{\beta})$ and $(\pi_{\gamma} | \delta_{\gamma})$ are updated by

$$\begin{aligned}p(\delta_{\beta_j} | \beta_j, \pi_{\beta_j}) &\propto p(\beta_j | \delta_{\beta_j}) p(\delta_{\beta_j} | \pi_{\beta_j}), \\ p(\delta_{\beta_j} = 1 | \beta_j, \pi_{\beta_j}) &= A / (A + B),\end{aligned}$$

where $A = \pi_{\beta_j} \phi(\beta_j; 0, s_j^2 \sigma_{\beta_j}^2)$, $B = (1 - \pi_{\beta_j}) \phi(\beta_j; \beta_j^p, \sigma_{\beta_j}^2)$. For $(\pi_{\beta_j} | \delta_{\beta_j})$,

$$\begin{aligned}p(\pi_{\beta_j} | \delta_{\beta_j}) &\propto p(\delta_{\beta_j} | \pi_{\beta_j}) p(\pi_{\beta_j}), \\ &= \pi_{\beta_j}^{\delta_{\beta_j}} (1 - \pi_{\beta_j})^{(1 - \delta_{\beta_j})} \pi_{\beta_j}^{1-1} (1 - \pi_{\beta_j})^{1-1}. \\ &\sim \text{Beta}(\delta_{\beta_j} + 1, 2 - \delta_{\beta_j}).\end{aligned}$$

For π_{γ_j} and δ_{γ_j} , the sampling method is similar.

To optimize mixing, we adjust the acceptance rate of the adaptive Metropolis-Hasting algorithm to around 0.2. Tuning parameters are increased(decreased) by multiplying(dividing) $\delta(l) = \exp(\min(0.01, l^{-1/2}))$ when current acceptance proportion is larger(smaller) than target optimal acceptance rate for every 100 iterations during burn-in period, where l is the number of current batches of 100 iterations (Jara et al., 2009). When the actual error distribution is far away from the Pólya tree baseline measure, the MCMC trace plot may reflect strong autocorrelation among posterior samples. Thus we recommend thinning to reduce the autocorrelation.

For the quantile regression coefficients, which are functionals of $(\beta, \gamma^*, \sigma, c, Y)$, we calculate the estimates by equation at each iteration using (9). Exact inference can be made using the posterior samples of the quantile regression coefficients (mean, median, and credible intervals).

3 Simulation Study

We conduct several simulation studies to compare our approach with other existing methods, specifically, the *rq* (RQ) function in the *quantreg* package (Koenker, 2012) in R Core Team (2013) (the standard frequentist quantile regression method) and the flexible Bayesian quantile regression approach by Reich (FBQR). We compare the approaches for both homogeneous and heterogeneous models.

3.1 Design

We generated data from the following 6 models,

$$\text{M1: } y_i = 1 + x_{i1}\beta_1 + \epsilon_{1i},$$

$$\text{M2: } y_i = 1 + x_{i1}\beta_1 + \epsilon_{2i},$$

$$\text{M3: } y_i = 1 + x_{i1}\beta_1 + \epsilon_{3i},$$

$$\text{M4: } y_i = 1 + x_{i1}\beta_1 + \epsilon_{4i},$$

$$\text{M1H: } y_i = 1 + x_{i1}\beta_1 + (1 + 0.2x_{i1})\epsilon_{1i},$$

$$\text{M2H: } y_i = 1 + x_{i1}\beta_1 + (1 + 0.2x_{i1})\epsilon_{2i},$$

$$\text{M3H: } y_i = 1 + x_{i1}\beta_1 + (1 + 0.2x_{i1})\epsilon_{3i},$$

$$\text{M4H: } y_i = 1 + x_{i1}\beta_1 + (1 + 0.2x_{i1})\epsilon_{4i},$$

$$\text{M5: } y_i | R_i = 1 \sim 2 + x_{i1} + \epsilon_{1i}, y_i | R_i = 0 \sim -2 - x_{i1} + \epsilon_{1i},$$

where $x_{i1} \stackrel{\text{iid}}{\sim} \text{Uniform}(0, 4)$, $\epsilon_{1i} \sim N(0, 1)$, $\epsilon_{2i} \sim t_3$, $\epsilon_{3i} \stackrel{\text{iid}}{\sim} 0.5 \times N(-2, 1) + 0.5 \times N(2, 1)$, $\epsilon_{4i} \sim 0.8N(0, 1) + 0.2N(3, 3)$. In model 1 (M1), the error distribution coincides with baseline distribution. Model 2 (M2) has a heavier tail distribution, student-t distribution with 3 degrees of freedom. Model 3 (M3) has a bimodal distribution for the error term. Model 4

(M4) uses a skewed mixture of normal distribution error introduced in [Reich et al. \(2010\)](#). Model 1H-4H (M1H-M4H) assume heterogeneous variances such that the quantiles lines are no longer parallel to each other. The heterogeneity in model 5 (M5) comes from the mixture of distributions not from covariates.

All covariates and error terms are mutually independent and all coefficients are set to be 1. For each model, we generate 100 data sets with the sample size $n = 200$. The quantiles estimated are 50% and 90%.

Each simulated data set is analyzed using the following four methods: RQ, FBQR, our proposed method with normal priors (PT), and our proposed method with spike and slab priors (PTSS). We used the default settings for RQ and FBQR. For PT, we adopt the following prior specifications:

$$\begin{aligned}\pi(\beta_j) &\sim N(\beta_j^p, V_{jj}), j = 0, 1, \\ \pi(\gamma_j) &\sim N(0, 100), j = 1, \\ \pi(\sigma) &\sim \Gamma(1/2, 1/2), \\ \pi(c) &\sim \Gamma(1/2, 1/2),\end{aligned}$$

where $\beta^p = (X'X)^{-1}X'y$ is the least square estimator, $V = \hat{\sigma}^2(X'X)^{-1}$, and $\hat{\sigma}^2 = \sum_{i=1}^n (y_i - x_i\beta^p)^2 / (n - 3)$. For PTSS, we used the same priors for σ and c , but spike-slab priors for β and γ :

$$\begin{aligned}\pi(\beta_j) &\sim \delta_{\beta_j}N(0, s_j V_{jj}) + (1 - \delta_{\beta_j})N(\beta_j^p, V_{jj}), j = 0, 1, \\ \pi(\gamma_j) &\sim \delta_{\gamma_j}N(0, 100s_j) + (1 - \delta_{\gamma_j})N(0, 100), j = 1, \\ \delta_{\beta_j} &\sim \text{Bernoulli}(\pi_{\beta_j}), \pi_{\beta_j} \sim \text{Beta}(1, 1), \\ \delta_{\gamma_j} &\sim \text{Bernoulli}(\pi_{\gamma_j}), \pi_{\gamma_j} \sim \text{Beta}(1, 1).\end{aligned}$$

And we choose $s_j = 1/1000$ from [George and McCulloch \(1993\)](#).

A partial Pólya tree with $M = 7$ levels was adopted. 180,000 iterations of a single Markov chain were used, during which, 30,000 samples were saved after a burn-in period of 30,000 samples by thinning every five samples for decreasing autocorrelation. It takes around 90 seconds for one simulation for PT under R version 2.15.3 (2013-03-01) and platform: x86_64-apple-darwin9.8.0/x86_64 (64-bit). Acceptance rates were set to approach 20% for all parameters candidates during the adaptive Metropolis-Hastings algorithm. We also tested the method proposed by Reich (FBQR), which conducts a single τ quantile regression for linear model and assigns an infinite mixture of Gaussian densities for the error term and the standard frequentist quantile regression approach, *rq* function in the *quantreg* package ([Koenker, 2012](#)) in [R Core Team \(2013\)](#) (RQ); both these approaches need a separate model/fit for each quantile of interest.

Methods are evaluated based on mean squared error:

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{\beta}_j(\tau) - \beta_j(\tau))^2,$$

where N is the number of simulations, $\beta_j(\tau)$ is the j^{th} component of the true quantile regression parameters. $\hat{\beta}_j(\tau)$ is the j^{th} component of estimated quantile regression parameters. We use the posterior mean as estimated parameters.

Monte Carlo standard errors (MCSE) are used to evaluate the “significance” of the differences between methods,

$$\text{MCSE} = \hat{\text{sd}}(\text{Bias}^2) / \sqrt{N},$$

where $\hat{\text{sd}}$ is the sample standard deviation and $\text{Bias} = \hat{\beta}_j(\tau) - \beta_j(\tau)$.

3.2 Results

The simulation results are shown in Table 1. In model 1 (M1), when the error is homogeneous and distributed as standard normal distribution, which coincides with Pólya tree baseline measure, RQ has a larger MSE than FBQR and PT. PTSS performed best since its prior shrunk the heterogeneity parameters toward zero. When considering heterogeneity in model 1 (M1H), PT and PTSS still perform well versus RQ and FBQR.

In model 2, 3, 4 and 2H, 3H, 4H, when error is homogeneous or heterogeneous, and is from a mixture of normals (student t distribution can be regarded as a mixture of normals), which is away from Pólya tree baseline measure, FBQR dominates the other three methods in terms of MSE, because simulated models coincide with the models assumed in the FBQR approach. However, PT and PTSS are also competitive. In median regression for model 3 and model 3 heterogeneity scenario, PT and PTSS have smaller MSE than FBQR and RQ. Similar situation happened in model 4 with 90% quantile regression.

In model 5, the heterogeneity comes from the mixture of distributions. The mode of the error distribution is no longer at median for RQ and FBQR, thus leading to larger MSE. Although PT and PTSS have larger MSE than RQ and BQR in 90% quantile, the deficit is offset by much smaller bias in 50% quantile regression.

To summarize, in all cases examined, the RQ method performs poorly in terms of MSE when the mode of the error is no longer the quantile of interest. In contrast, PT is not impacted by lack of unimodality and heterogeneity and provides more information for the relationship between responses and covariates; both RQ and FBQR perform less well than PT/PTSS when the quantile of interest is not near the mode of the error. FBQR outperforms PT in some cases due to the flexibility of the distribution of the error being assigned an infinite mixture of normal distribution. Less information is available from our approach to detect the shape at a particular extreme percentile of the distribution since there are few observations at extreme quantiles. However, PT and PTSS can fit simultaneously multiple quantile regressions and provide coherent information about the error distribution. An overall evaluation method over multiple quantiles, such as summation of MSE over all quantiles and coefficients, may reflect PT and PTSS have advantages when error distribution is away from regular unimodal shape as in model 3 (M3 and M3H) and model 5 (M5). Meanwhile, quantile lines do not cross using our method, while we observed crossing quantile lines from FBQR method for several of the simulated datasets. We also expect to see advantages when dimension of responses is larger than one.

4 Data Analysis

In this section, we apply our Bayesian quantile regression approach to examine the quantiles of 6 month weight loss from a recent weight management study, TOURS (Perri et al., 2008). This trial was designed to test whether a lifestyle modification program could effectively help people to manage their weights in the long term. In particular, we are interested in the effects of age and race. We focus on the weight loss from baseline to 6 months. The age of the subjects ranged from 50 to 75, and there were 43 people with race classified as black and 181 people as white. Our goal is to determine how the percentiles of weight change are affected by their age and race. “Age” covariate are scaled to 0 to 5 with every increment representing 5 years.

We fitted regression models for quantiles (10%, 30%, 50%, 70%, 90%). And we used Bayesian posterior samples to construct 95% credible intervals.




Figure 1: Scatterplots of weight loss vs age and Boxplots of weight loss for each race. The boxplots use the default settings: (0.75, 0.5, 0.25) quantile for box and $Q1 - 1.5IQR$ for lower whisker and $Q3 + 1.5IQR$ for upper whisker.

Results appear in Table 2. Whites lost more weight than blacks for all quantiles. The differential is reported as significant and becomes larger when comparing more successful weight losers (70% - 90% percentile). For example, whites lost 5.05 kg more than blacks among people losing the most weight (90%) reported from method PT (4.21 kg from PTSS).

The effect of age on the weight loss is small and not significant in most cases (only barely significant in 10% and 30% quantile regression by PT and PTSS). The trend is negative showing that older people tend to lose less weight. For example, median weight loss is 0.40 kg less for every age increase of 5 years reported by PT.

PTSS tends to shrink coefficients toward zero. For example, the posterior probability that the heterogeneity parameters are zero are 1.0 for age and .99 for race, indicating there is no heterogeneity for covariates for age and race. These priors can help to select variables and increase efficiency in Bayesian models. For example, we can exclude AGE out of the regressors and conclude the variance is homogeneous on the AGE covariate.

5 Discussion

This paper introduced a Bayesian approach for simultaneous linear quantile regression by introducing mixture of Pólya tree priors and estimating heterogeneity parameters. By marginalizing the predictive density function of the Pólya tree distribution, quantiles of interest can be obtained in closed form by inverting the predictive cumulative distribution. Exact posterior inference can be made via MCMC. In addition, traditional frequentist method and FBQR show crossing quantile lines for both simulated and real data. However, quantile lines from our method cannot cross since quantiles are estimated through density estimation. The simulations show advantages of our method in some cases especially when the error is multimodal and highly skewed. For multiple quantiles of interest, our method only needs one model and avoids crossing quantile curves, while FBQR and RQ need separate models assuming different error distributions for each quantile.

Further research includes quantile regression for correlated data by modelling error as a mixture of multivariate Pólya tree distribution. Our approach allows for quantile regression with missing data under ignorability by adding a data augmentation step. We are exploring extending our approach to allow for nonignorable missingness. Also it might be possible to use a slightly more complex baseline distribution in Pólya tree adaptively to improve the estimation.

References

- J.O. Berger and A. Guglielmi. Bayesian and conditional frequentist testing of a parametric model versus nonparametric alternatives. *Journal of the American Statistical Association*, 96 (453):174–184, 2001.
- Moshe Buchinsky. Changes in the US Wage Structure 1963-1987: Application of Quantile Regression. *Econometrica*, 62(2):pp. 405–458, 1994. ISSN 00129682. URL <http://www.jstor.org/stable/2951618>.
- Moshe Buchinsky. Recent advances in quantile regression models: A practical guideline for empirical research. *The Journal of Human Resources*, 33(1):pp. 88–126, 1998. ISSN 0022166X. URL <http://www.jstor.org/stable/146316>.
- J. Fabius. Asymptotic behavior of bayes’ estimates. *The Annals of Mathematical Statistics*, 35 (2):846–856, 1964.
- T.S. Ferguson. Prior distributions on spaces of probability measures. *The Annals of Statistics*, pages 615–629, 1974.
- D.A. Freedman. On the asymptotic behavior of bayes’ estimates in the discrete case. *The Annals of Mathematical Statistics*, 34(4):1386–1403, 1963.
- Edward I. George and Robert E. McCulloch. Variable selection via gibbs sampling. *Journal of the American Statistical Association*, 88(423):pp. 881–889, 1993. ISSN 01621459. URL <http://www.jstor.org/stable/2290777>.

- J. Geweke and M. Keane. Smoothly mixing regressions. *Journal of Econometrics*, 138(1):252–290, 2007.
- T. Hanson and W.O. Johnson. Modeling regression error with a mixture of polya trees. *Journal of the American Statistical Association*, 97(460):1020–1033, 2002.
- X. He, P. Ng, and S. Portnoy. Bivariate quantile smoothing splines. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 60(3):537–550, 1998. ISSN 1467-9868. doi: 10.1111/1467-9868.00138. URL <http://dx.doi.org/10.1111/1467-9868.00138>.
- N.L. Hjort and S. Petrone. Nonparametric quantile inference using dirichlet processes. *Advances in statistical modeling and inference*, pages 463–492, 2007.
- N.L. Hjort and S.G. Walker. Quantile pyramids for bayesian nonparametrics. *The Annals of Statistics*, 37(1):105–131, 2009.
- A. Jara, T.E. Hanson, and E. Lesaffre. Robustifying generalized linear mixed models using a new class of mixtures of multivariate polya trees. *Journal of Computational and Graphical Statistics*, 18(4):838–860, 2009.
- Alejandro Jara, Timothy Hanson, Fernando Quintana, Peter Müller, and Gary Rosner. DP-package: Bayesian semi- and nonparametric modeling in R. *Journal of Statistical Software*, 40(5):1–30, 2011. URL <http://www.jstatsoft.org/v40/i05/>.
- R. Koenker. *Quantile regression*, volume 38. Cambridge Univ Pr, 2005.
- Roger Koenker. *quantreg: Quantile Regression*, 2012. URL <http://CRAN.R-project.org/package=quantreg>. R package version 4.91.
- Roger Koenker and Jr. Bassett, Gilbert. Regression quantiles. *Econometrica*, 46(1):pp. 33–50, 1978. ISSN 00129682. URL <http://www.jstor.org/stable/1913643>.
- Roger Koenker and Jose A. F. Machado. Goodness of fit and related inference processes for quantile regression. *Journal of the American Statistical Association*, 94(448):pp. 1296–1310, 1999. ISSN 01621459. URL <http://www.jstor.org/stable/2669943>.
- A. Kottas and A.E. Gelfand. Bayesian semiparametric median regression modeling. *Journal of the American Statistical Association*, 96(456):1458–1468, 2001.
- A. Kottas and M. Krnjajić. Bayesian semiparametric modelling in quantile regression. *Scandinavian Journal of Statistics*, 36(2):297–319, 2009.
- M. Lavine. Some aspects of polya tree distributions for statistical modelling. *The Annals of Statistics*, pages 1222–1235, 1992.
- M. Lavine. More aspects of polya tree distributions for statistical modelling. *The Annals of Statistics*, pages 1161–1176, 1994.
- R.D. Mauldin, W.D. Sudderth, and SC Williams. Polya trees and random distributions. *The Annals of Statistics*, pages 1203–1221, 1992.

- S.M. Paddock. *Randomized Polya trees: Bayesian nonparametrics for multivariate data analysis*. PhD thesis, Duke University, 1999.
- Michael G Perri, Marian C Limacher, Patricia E Durning, David M Janicke, Lesley D Lutes, Linda B Bobroff, Martha Sue Dale, Michael J Daniels, Tiffany A Radcliff, and A Daniel Martin. Extended-care programs for weight management in rural communities: the treatment of obesity in underserved rural settings (tours) randomized trial. *Archives of internal medicine*, 168(21):2347, 2008.
- R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2013. URL <http://www.R-project.org/>. ISBN 3-900051-07-0.
- B.J. Reich, H.D. Bondell, and H.J. Wang. Flexible bayesian quantile regression for independent and clustered data. *Biostatistics*, 11(2):337–352, 2010.
- Brian J. Reich. Spatiotemporal quantile regression for detecting distributional changes in environmental processes. *J. R. Stat. Soc. Ser. C. Appl. Stat.*, 61(4):535–553, 2012. ISSN 0035-9254. doi: 10.1111/j.1467-9876.2011.01025.x. URL <http://dx.doi.org/10.1111/j.1467-9876.2011.01025.x>.
- Brian J. Reich and Luke B. Smith. Bayesian quantile regression for censored data. *Biometrics*, 69(3):651–660, 2013. ISSN 0006-341X. doi: 10.1111/biom.12053. URL <http://dx.doi.org/10.1111/biom.12053>.
- Brian J. Reich, Montserrat Fuentes, and David B. Dunson. Bayesian spatial quantile regression. *J. Amer. Statist. Assoc.*, 106(493):6–20, 2011. ISSN 0162-1459. doi: 10.1198/jasa.2010.ap09237. URL <http://dx.doi.org/10.1198/jasa.2010.ap09237>.
- L. Scaccia and P.J. Green. Bayesian growth curves using normal mixtures with nonparametric weights. *Journal of Computational and Graphical Statistics*, 12(2):308–331, 2003.
- M.A. Taddy and A. Kottas. A bayesian nonparametric approach to inference for quantile regression. *Journal of Business and Economic Statistics*, 28(3):357–369, 2010.
- S. Tokdar and J.B. Kadane. Simultaneous linear quantile regression: A semiparametric bayesian approach. *Bayesian Analysis*, 6(4):1–22, 2011.
- S.G. Walker and B.K. Mallick. Hierarchical generalized linear models and frailty models with bayesian nonparametric mixing. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 59(4):845–860, 1997.
- Stephen Walker and Bani K. Mallick. A bayesian semiparametric accelerated failure time model. *Biometrics*, 55(2):477–483, 1999. ISSN 1541-0420. doi: 10.1111/j.0006-341X.1999.00477.x. URL <http://dx.doi.org/10.1111/j.0006-341X.1999.00477.x>.
- Ying Wei, Anneli Pere, Roger Koenker, and Xuming He. Quantile regression methods for reference growth charts. *Statistics in Medicine*, 25(8):1369–1382, 2006. ISSN 1097-0258. doi: 10.1002/sim.2271. URL <http://dx.doi.org/10.1002/sim.2271>.

Keming Yu and Rana A Moyeed. Bayesian quantile regression. *Statistics & Probability Letters*, 54(4):437–447, 2001.

Keming Yu, Zudi Lu, and Julian Stander. Quantile regression: applications and current research areas. *Journal of the Royal Statistical Society: Series D (The Statistician)*, 52(3):331–350, 2003. ISSN 1467-9884. doi: 10.1111/1467-9884.00363. URL <http://dx.doi.org/10.1111/1467-9884.00363>.

Table 1: Mean squared error (reported as 100*average) and MCSE (reported as 100*MCSE) for each quantile regression method. The four columns (RQ, FBQR, PT, PTSS) stand for frequentist method *rq* function from *quantreg* R package (RQ), flexible Bayesian method by Reich (FBQR), and our Bayesian approach using Pólya tree with normal priors (PT) and with spike and slab priors (PTSS).

	RQ	FBQR	PT	PTSS	RQ	FBQR	PT	PTSS
	M1 50%				M1H 50%			
β_0	2.55(0.39)	1.69(0.23)	1.70(0.23)	1.70(0.23)	3.05(0.60)	2.38(0.42)	2.41(0.40)	2.42(0.39)
β_1	0.52(0.08)	0.31(0.04)	0.31(0.04)	0.31(0.04)	0.84(0.18)	0.54(0.11)	0.60(0.11)	0.60(0.11)
	M1 90%				M1H 90%			
β_0	7.68(0.98)	4.89(0.70)	3.91(0.54)	3.39(0.48)	9.53(1.82)	6.90(12.9)	4.80(1.02)	5.76(1.02)
β_1	1.31(0.16)	0.84(0.12)	0.73(0.10)	0.60(0.08)	2.33(0.41)	1.60(0.43)	1.33(0.26)	1.49(0.27)
	M2 50%				M2H 50%			
β_0	3.41(0.41)	2.67(0.37)	2.83(0.38)	2.77(0.36)	4.23(0.53)	2.84(0.31)	4.54(0.54)	4.92(0.58)
β_1	0.61(0.07)	0.49(0.06)	0.51(0.07)	0.51(0.06)	0.96(0.15)	0.73(0.10)	1.23(0.17)	1.30(0.17)
	M2 90%				M2H 90%			
β_0	18.1(3.31)	12.0(2.01)	16.1(2.72)	11.9(2.27)	32.8(6.52)	15.1(2.40)	20.7(4.04)	30.9(4.37)
β_1	3.64(0.57)	1.84(0.22)	3.39(0.47)	2.00(0.31)	8.35(1.30)	3.70(0.62)	7.95(1.43)	5.62(1.20)
	M3 50%				M3H 50%			
β_0	82.0(8.61)	16.6(2.36)	9.02(1.40)	13.7(2.29)	98.6(9.87)	16.5(2.09)	10.3(1.29)	16.3(2.12)
β_1	17.7(1.73)	1.79(0.33)	1.21(0.20)	1.19(0.19)	26.9(2.98)	2.88(0.36)	2.06(0.29)	1.80(0.25)
	M3 90%				M3H 90%			
β_0	10.9(1.34)	6.64(0.99)	9.26(1.53)	8.99(1.45)	13.2(1.89)	9.11(1.28)	12.3(1.70)	12.4(1.75)
β_1	2.16(0.30)	1.41(0.17)	1.76(0.26)	1.68(0.24)	3.94(0.51)	2.61(0.36)	3.58(0.50)	3.56(0.51)
	M4 50%				M4H 50%			
β_0	5.74(0.75)	4.26(0.61)	6.47(0.87)	6.46(0.90)	5.24(0.74)	5.09(0.65)	6.66(0.83)	7.06(0.95)
β_1	0.84(0.09)	0.61(0.08)	0.86(0.12)	0.84(0.12)	1.42(0.19)	1.14(0.15)	1.38(0.18)	1.44(0.18)
	M4 90%				M4H 90%			
β_0	53.0(6.46)	23.2(3.08)	23.0(2.88)	19.6(2.52)	88.7(11.1)	37.1(4.88)	38.5(5.02)	40.5(5.67)
β_1	11.4(1.39)	3.10(0.36)	3.83(0.48)	2.79(0.36)	25.0(3.17)	7.09(0.94)	7.46(0.94)	8.04(1.00)
	M5 50%							
β_0	174(18.6)	46.0(5.92)	6.77(1.44)	4.87(1.06)				
β_1	150(10.7)	10.8(1.68)	1.83(0.61)	1.58(0.33)				
	M5 90%							
β_0	8.10(1.07)	6.68(0.97)	11.5(1.48)	13.1(1.69)				
β_1	1.54(0.22)	2.03(0.28)	2.77(0.36)	2.96(0.40)				

Table 2: 95% credible (confidence) intervals for quantile regression parameters for TOURS. PT is our proposed Pólya trees approach with normal priors, and PTSS, Pólya trees approach with spike-slab priors .

Term	PT	PTSS
10%		
Intercept	2.62(1.11,4.22)	2.10(0.65,3.36)
Age	-0.57(-1.25,-0.03)	-0.57(-1.09,-0.07)
Race	2.70(1.20,4.29)	3.32(2.07,4.70)
30%		
Intercept	5.59(4.64,6.70)	5.45(4.41,6.36)
Age	-0.46(-0.91,-0.10)	-0.47(-0.82,-0.19)
Race	3.38(2.22,4.42)	3.58(2.56,4.65)
50%		
Intercept	7.43(6.46,8.56)	7.47(6.24,8.40)
Age	-0.40(-0.75,-0.08)	-0.42(-0.72,-0.16)
Race	3.81(2.77,4.68)	3.74(2.76,4.72)
70%		
Intercept	9.79(8.74,11.1)	10.1(8.92,11.2)
Age	-0.31(-0.74,0.06)	-0.34(-0.74,0.00)
Race	4.35(3.19,5.39)	3.94(2.87,4.99)
90%		
Intercept	12.8(11.3,14.6)	13.5(12.0,15.1)
Age	-0.20(-0.89,0.38)	-0.24(-0.86,0.30)
Race	5.05(3.36,6.61)	4.21(2.85,5.51)